

AI Accelerator Design Automation: From Network to RTL

Team members: Wen-Cong Huang, Zih-Sing Fu

Advisor: Prof. Chia-Hsiang Yang

Technical Highlights

❖ A flexible AI accelerator

An engine for flexible network inference, including highly configurable descriptor, ISA and hardware co-design

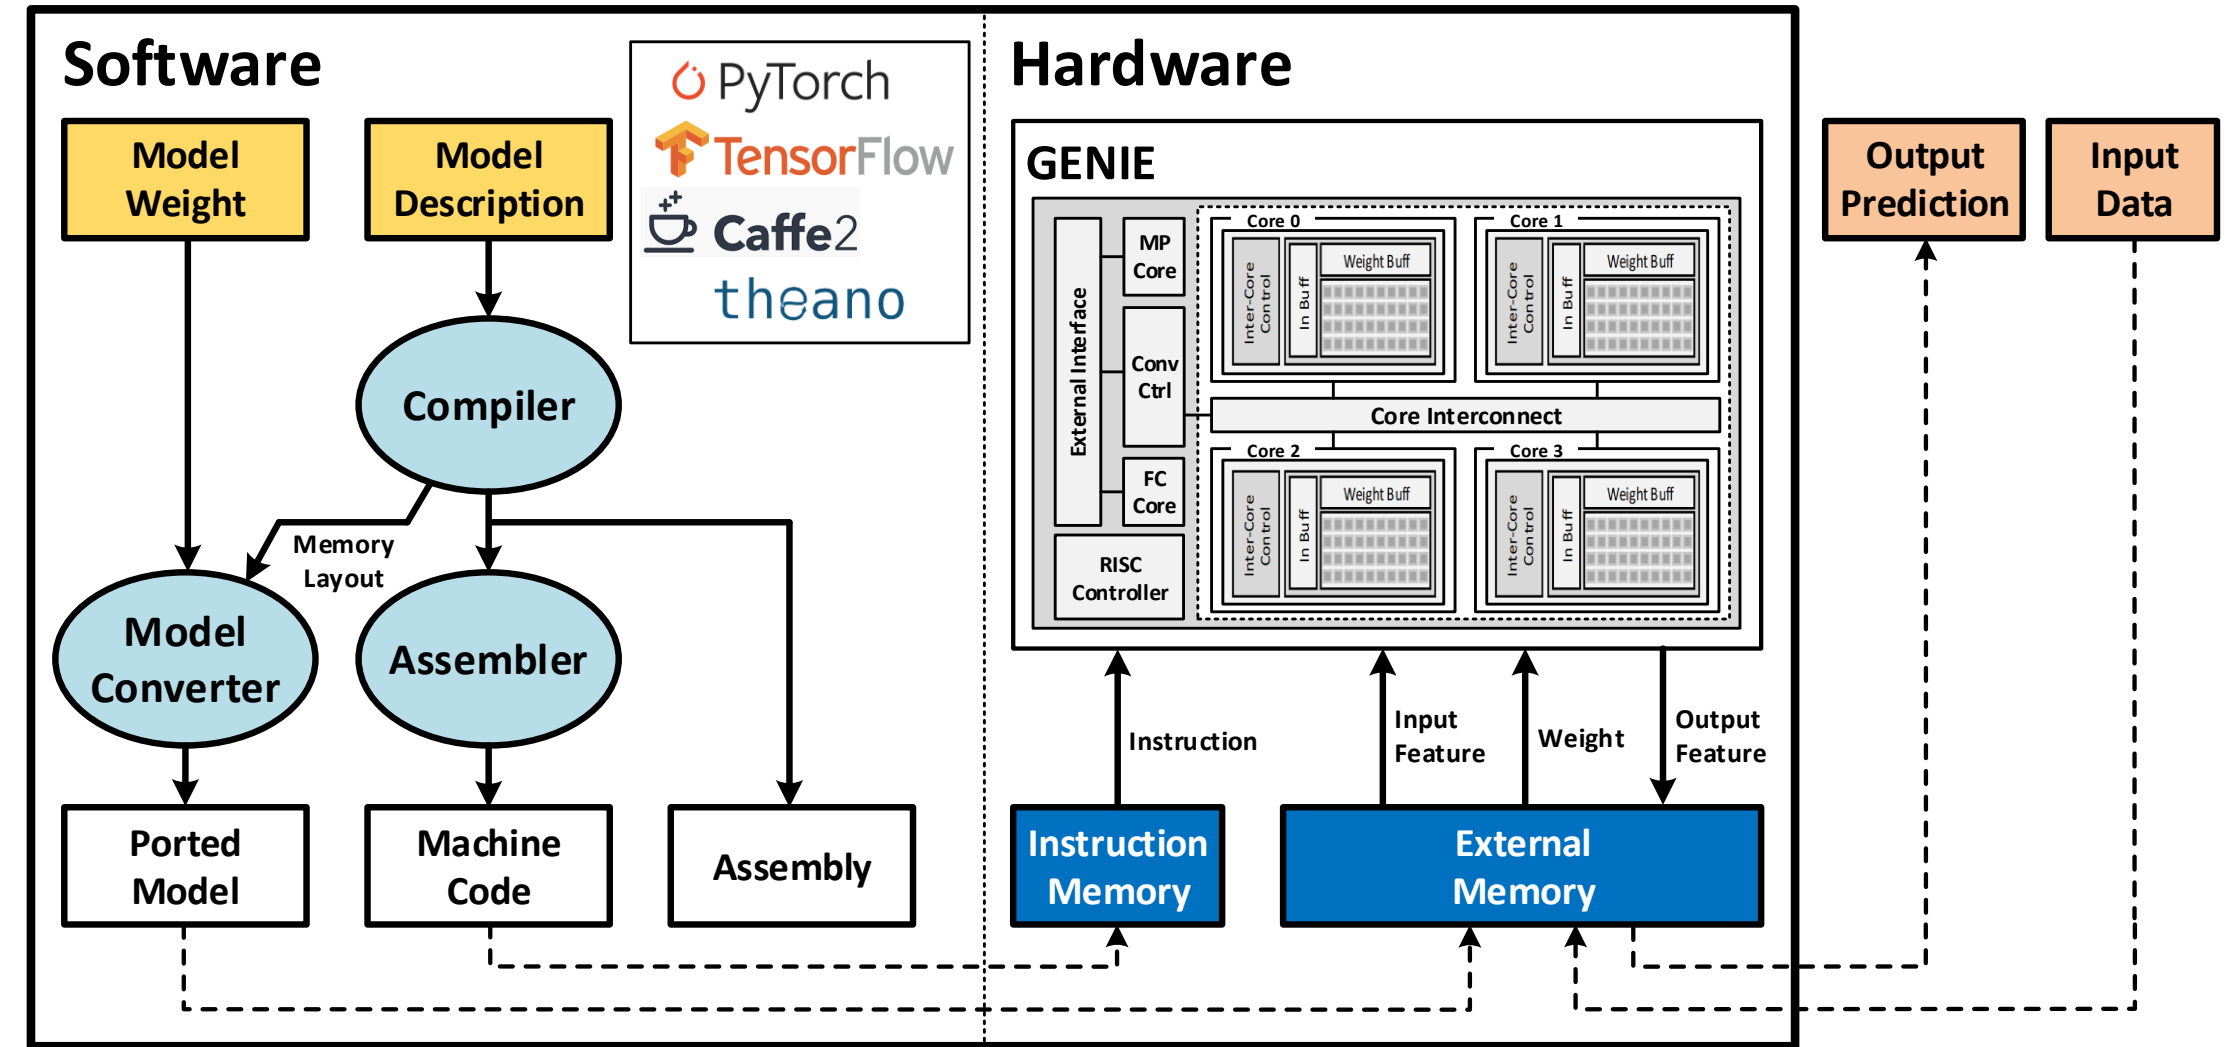
❖ An easy-to-use AI accelerator

Support complete inference flow of arbitrary network architecture with provided model descriptor and weight

❖ Technical focus

- ❑ Model compiler from custom model descriptor
- ❑ Software and hardware co-design for AI applications
- ❑ Optimal data scheduling for multi-core processing
- ❑ Optimal tensor partition strategy exploration

Design Automation and User-friendly Workflow



Flexible Model Descriptor

❖ Sequential model descriptor

❖ Supported operations

- ❑ 2D convolutional layer
 - Arbitrary input feature shape
 - Flexible activation/bias toggle
- ❑ Fully-connected layer
- ❑ Max pooling layer
- ❑ ReLU activation

```
model = nn.Sequential(  
    nn.Conv2d(1, 16, 3, padding=1, bias=True),  
    nn.Conv2d(16, 16, 3, padding=1, bias=True),  
    nn.Conv2d(16, 16, 3, padding=1, bias=True),  
    nn.ReLU(inplace=True),  
    nn.MaxPool2d(2),  
    nn.Conv2d(16, 32, 3, padding=1, bias=False),  
    nn.Conv2d(32, 32, 3, padding=1, bias=False),  
    nn.Conv2d(32, 32, 3, padding=1, bias=False),  
    nn.ReLU(inplace=True),  
    nn.MaxPool2d(2),  
    nn.Conv2d(32, 64, 3, padding=1, bias=True),  
    nn.Conv2d(64, 64, 3, padding=1, bias=True),  
    nn.Conv2d(64, 64, 3, padding=1, bias=True),  
    nn.ReLU(inplace=True),  
    nn.MaxPool2d(2),  
    Flatten(),  
    nn.Linear(576, 10, bias=False),  
)
```

```
1 ifdim 1,28,28  
2 conv 1,16,3,1 noact bias  
3 conv 16,16,3,1 noact bias  
4 conv 16,16,3,1 relu bias  
5 maxpool  
6 conv 16,32,3,1 noact nobias  
7 conv 32,32,3,1 noact nobias  
8 conv 32,32,3,1 relu nobias  
9 maxpool  
10 conv 32,64,3,1 noact bias  
11 conv 64,64,3,1 noact bias  
12 conv 64,64,3,1 relu bias  
13 maxpool  
14 flatten  
15 fc 576,10 relu nobias
```

model
descriptor

compiled
instructions

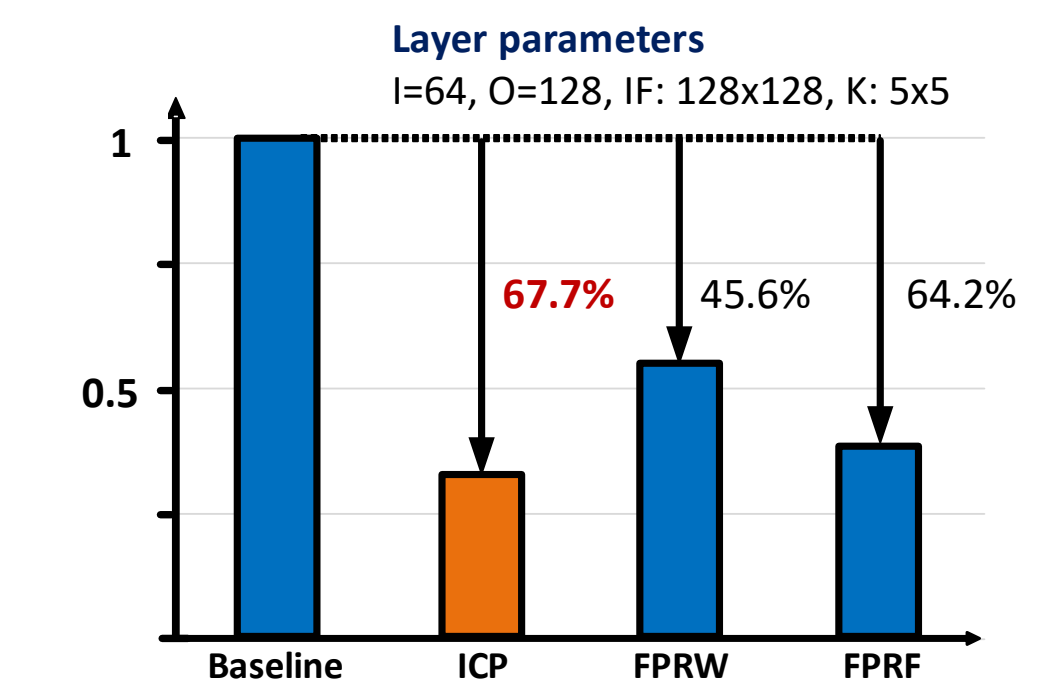
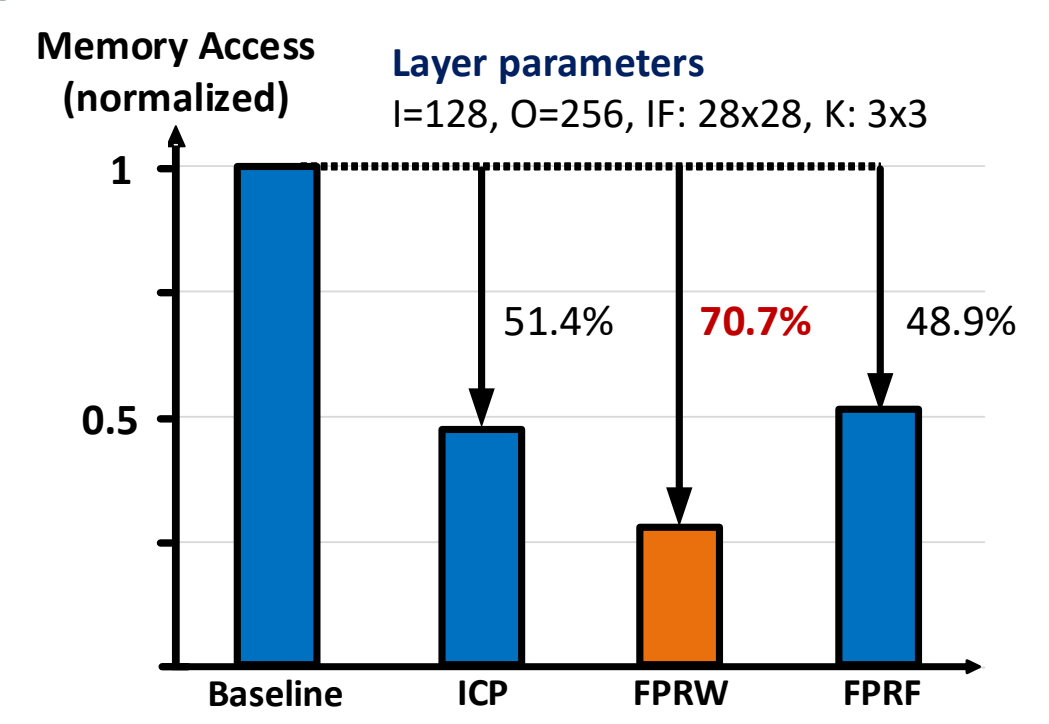
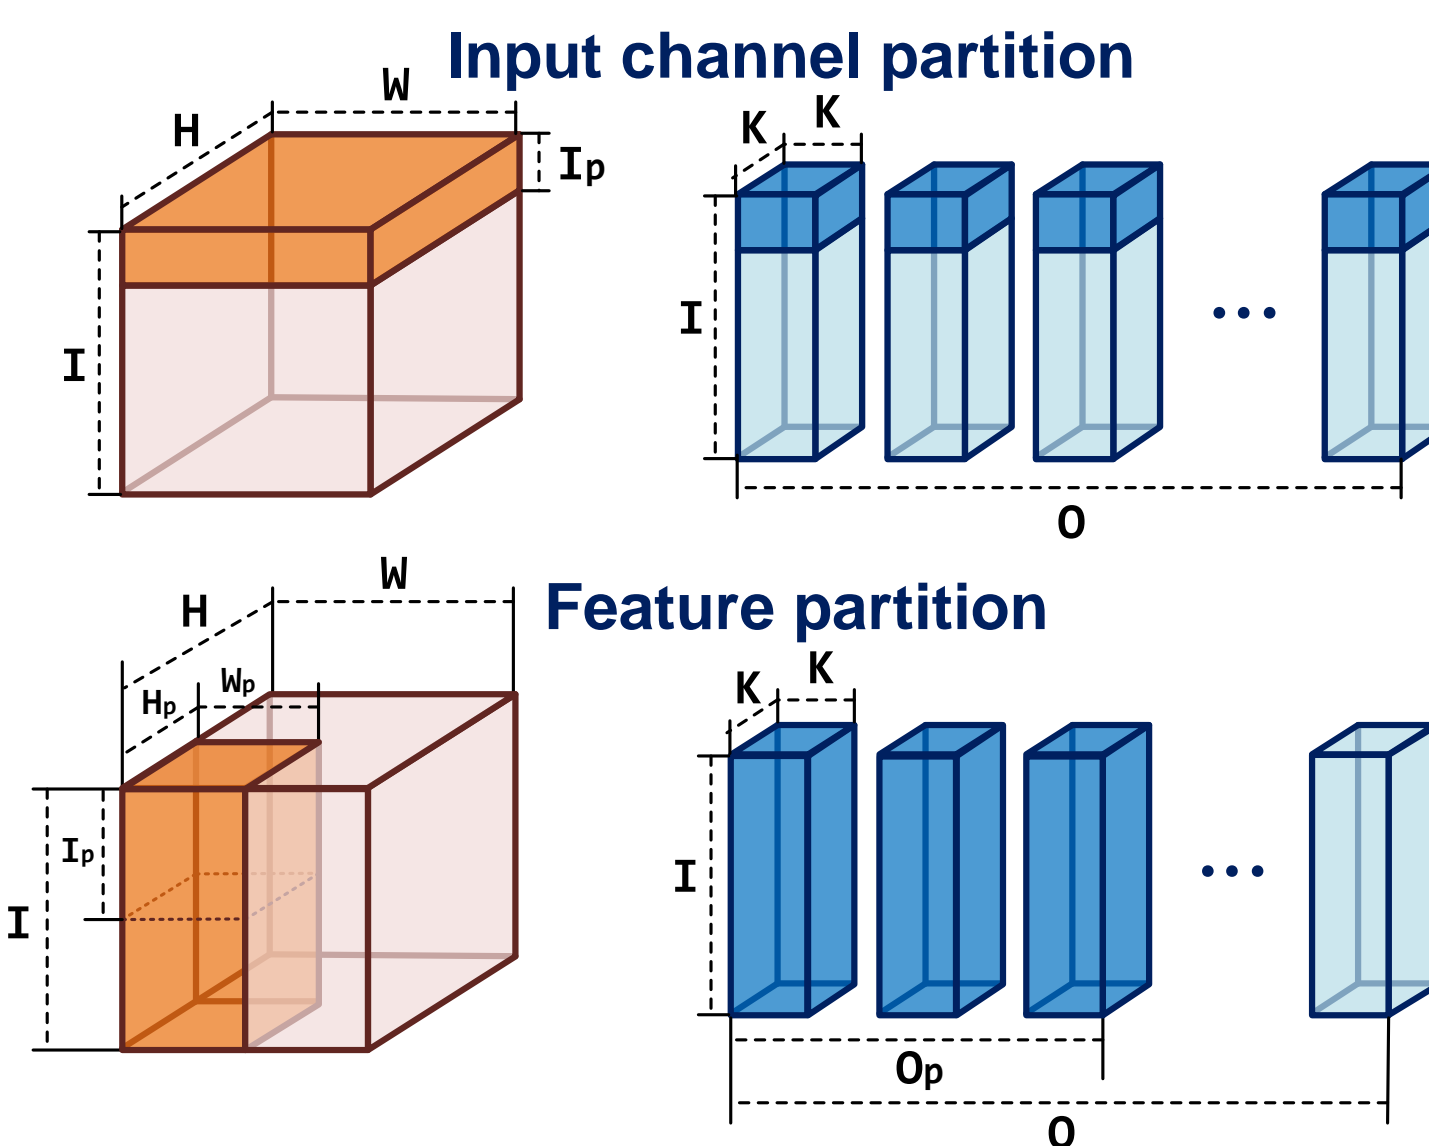
Memory Scheduling and Optimization

❖ Support for arbitrary feature size require tensor partitioning

❖ Memory access and on-chip memory should be minimized

❖ Data scheduling strategies

1. **Baseline** (no partition reuse)
2. **ICP**: Input channel partition
3. **FPRW**: Feature partition (reload weight)
4. **FPRF**: Feature partition (reload feature)



The compiler selects different strategies to optimize memory access for different layers

ISA and Instruction Generation

❖ Instruction Set Architecture

- ❑ 32-bit RISC ISA
- ❑ 19 instructions

❖ Instruction generation

Algorithm 6 Convolution Instruction Generation Routine (Single Core)

```
1: procedure CONVROUTINE  
2:   cfgl conv, activation, bias  
3:   cfgecv I, O, K, pad  
4:   cfgecvif H, W  
5:   cvaif ifbase  
6:   cvaw wbase  
7:   cvaof ofbase  
8:   for  $O_{ori} \leftarrow 0$  to  $O$  step  $O_p$  do  
9:     cvcfgori 0,  $O_p$   
10:    cvcfgevt 0,  $\min(O - O_p, O_p)$   
11:    cvlwp  
12:    for  $H_{ori} \leftarrow -pad$  to  $H + pad$  step  $H_p - K + 1$  do  
13:      for  $W_{ori} \leftarrow -pad$  to  $W + pad$  step  $W_p - K + 1$  do  
14:        cvcfgori  $H_{ori}, W_{ori}$   
15:        cvcfgevt  $\min(H + pad - H_{ori}, H_p), \min(W + pad - W_{ori}, W_p)$   
16:        for  $I_{ori} \leftarrow 0$  to  $I$  step  $I_p$  do  
17:          cvcfgori  $I_{ori}, hold$   
18:          cvcfgevt  $\min(I - I_{ori}, I_p), hold$   
19:          cvlifp  
20:          cvsofp
```

Configurations

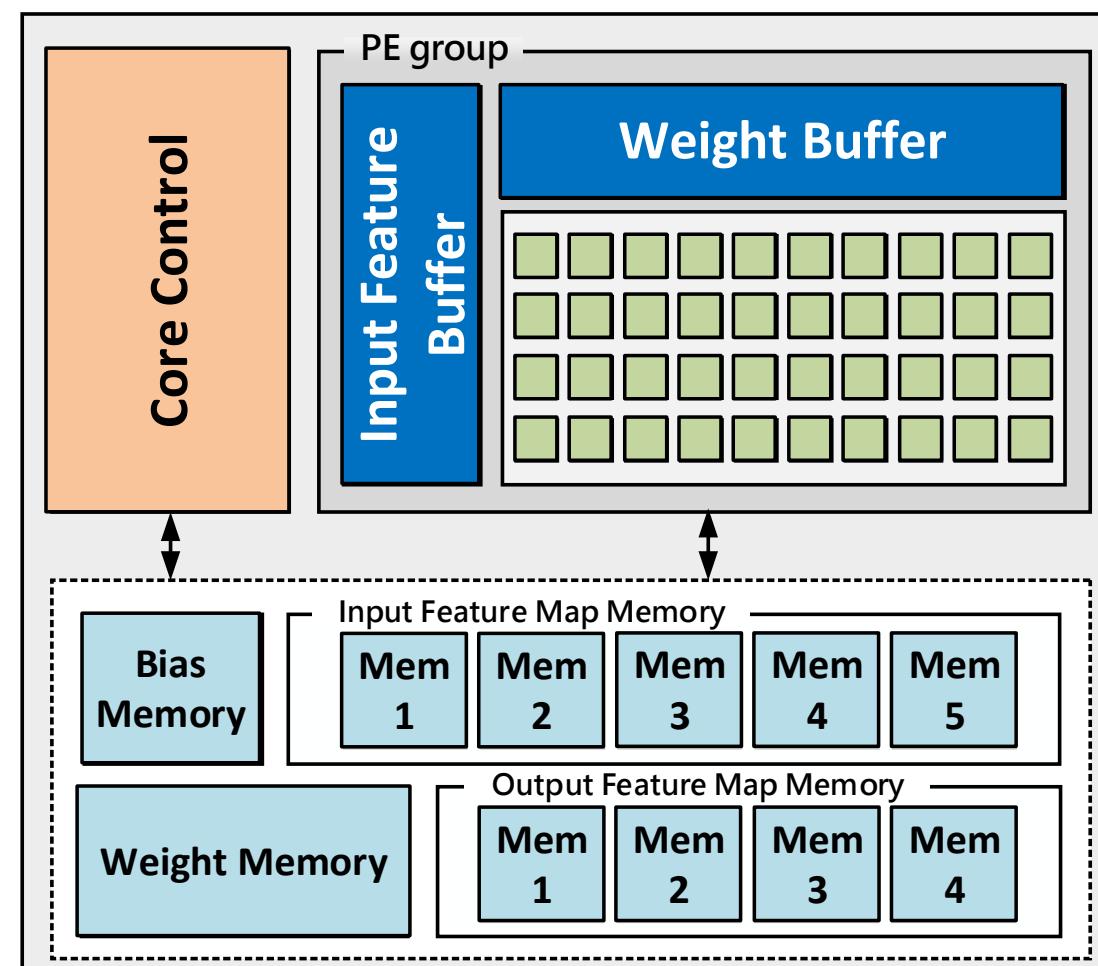
cfgecv	31:27	26:16	15:5	4:3	2:0
cfgecv	11	Cin	Cout	padding	kernel

Multi-core support

cvseipe	31:27	26:9	8	7:0
cvseipe	16	reserved	broadcast	peid

AI Accelerator Design

❖ Core Architecture



❖ System Specifications

- ❑ 4-core architecture for bandwidth improvement
- ❑ Configurable memory options for performance tweaking

❖ Memory Estimation

Memory Type	Memory size	Estimated area
Ifmap Memory	~25 KB	$175000 \mu m^2$
Ofmap Memory	~18.5 KB	$129500 \mu m^2$
Weight Memory	~73.7 KB	$515900 \mu m^2$
Bias Memory	~128B	$896 \mu m^2$
Total Memory	~120KB	$840000 \mu m^2$

Research Outcomes

- ❖ **Practical**: from network model to RTL in minutes
- ❖ **Efficient**: more than 60% less external memory access
- ❖ **Reconfigurable**: different memory size for different scenarios
- ❖ **Ongoing tasks**: AXI-compatibility, higher external throughput